

## Problem Set 4, Part I

### Problem 1: Timestamp-based concurrency control

1.1

request	response of the system	any changes to state for A <i>and/or</i> explanation of why action wasn't accepted
r2(A)	allowed	RTS = 20
w4(A)	allowed	WTS = 40
r3(A)	denied; rollback	$TS(T3) < WTS(A)$ ; $30 < 40$
w1(A)	Denied; rollback	$TS(T1) < RTS(A)$ ; $10 < 20$
w2(A)	ignored	$TS(T2) \geq RTS(A)$ but $TS(T2) < WTS(A)$
r5(A)	allowed	RTS = 50

1.2

request	response of the system	any changes to state for A <i>and/or</i> explanation of why action wasn't accepted
r2(A)	allowed	RTS = 20
w4(A)	allowed	WTS = 40; c = false
r3(A)	denied; rollback	$TS(T3) < WTS(A)$ ; $30 < 40$ ; c = false
w1(A)	Denied; rollback	$TS(T1) < RTS(A)$ ; $10 < 20$ ; c = false
w2(A)	Denied; make wait	$TS(T2) \geq RTS(A)$ but $TS(T2) < WTS(A)$ but c = false
r5(A)	allowed	RTS = 50
c2	allowed	
c4	allowed	c = true, WTS = 40
c5	allowed	
w2(A)	ignored	Already committed

1.3

request	response of the system	any changes to state for A <i>and/or</i> explanation of why action wasn't accepted
r2(A)	allowed to read A(0)	$RTS(A(0)) = 20$ , $WTS(A(0)) = 0$
w4(A)	Allowed to write A(40)	$WTS(A(40)) = 40$ , <b><math>RTS(A(40)) = 0</math></b>
r3(A)	Allowed to read A(0)	Can read from A(0), A(40), A(0) is valid; <b><math>RTS(A(0)) = 30</math></b>
w1(A)	denied	No such value of A where <b><math>TS &gt; WTS(A)</math></b>
w2(A)	Allowed to write A(0)	$WTS(A(0)) = 20$ ; [largest <b><math>WTS(A) \leq TS</math></b> is 20, allowed because read time stamp is equal]
r5(A)	Allowed to read A(40)	<b><math>RTS(A(40)) = 50</math></b>

## **Problem 2: Replication and distributed concurrency control**

2.1) 11 copies of each item, fully distributed locking

<b><i>voting scheme</i></b>	<b><i>would it work? (yes/no)</i></b>	<b><i>explanation</i></b>
1a) update 6, read 5	no	Although we update a majority of copies, we might need 5 stale ones. $r$ is not $> n - w$
1b) update 5, read 8	no	<b>Not updating a majority of the copies</b> , although the read condition is satisfied
1c) update 8, read 4	yes	<b>Updating majority</b> ( $w > n/2$ ) and also reading $> n-w$ because $4 > 11-8$
1d) update 3, read 10	no	<b>Not updating majority of copies</b>

2.2) 11 copies of each item, primary-copy locking

<b><i>voting scheme</i></b>	<b><i>would it work? (yes/no)</i></b>	<b><i>explanation</i></b>
1a) update 6, read 5	no	$r$ is not $> n-w$ since $5 \neq 11 - 6$
1b) update 5, read 8	yes	$r > n - w \Rightarrow 8 > 11-5=6$ , read majority of copies
1c) update 8, read 4	yes	$r > n - w \Rightarrow 4 > 11-8 = 3$ , read majority
1d) update 3, read 10	yes	$r > n - w \Rightarrow 10 > 11-3 = 8$ , read majority

2.3)

Using a configuration with more reads can guarantee fault tolerance. **1) (d) would be a good choice** because we have **10 reads out of 11, and 3 updates**. In a read-heavy workload, we ensure **not to read any stale data** because  $r > n-w$  and even if we have a site go down, we have a **very high probability of reading an updated value without incurring latency costs**. Networking would not be too much of a bottleneck either since the data can be read from 10 locations.

2.4)

For a write heavy workload, although it may feel like **1 (c) is a good choice** because we have 8 writers, writes usually incur a higher latency cost. It would seem wise to use **1 (a)** because you have a majority of writers which can deal with failed machines and still ensure consistency, and at the same time ensure we **do not incur too high of a writing cost**.