

Problem Set 4, Part I

Problem 1: Timestamp-based concurrency control

1.1

request	response of the system	any changes to state for A <i>and/or</i> explanation of why action wasn't accepted
r2(A)	allowed	RTS = 20
w4(A)	allowed	WTS = 40
r3(A)	denied; rollback	$TS(T3) < WTS(A)$; $30 < 40$
w1(A)	Denied; rollback	$TS(T1) < RTS(A)$; $10 < 20$
w2(A)	ignored	$TS(T2) \geq RTS(A)$ but $TS(T2) < WTS(A)$
r5(A)	allowed	RTS = 50

1.2

request	response of the system	any changes to state for A <i>and/or</i> explanation of why action wasn't accepted
r2(A)	allowed	RTS = 20
w4(A)	allowed	WTS = 40; $c(A) = \text{false}$
r3(A)	denied; rollback	$TS(T3) < WTS(A)$; $30 < 40$; $c(A) = \text{false}$
w1(A)	Denied; rollback	$TS(T1) < RTS(A)$; $10 < 20$; $c(A) = \text{false}$
w2(A)	Denied; make wait	$TS(T2) \geq RTS(A)$ ($20 \geq 20$) but $TS(T2) < WTS(A)$ [$20 < 40$] but $c(A) = \text{false}$
r5(A)	Denied make wait	$RTS = 50 \geq WTS(A)$ but $c = \text{false}$
c1	denied	
c3	denied	
c4	allowed	$c(A) = \text{true}$, WTS = 40
w2(A)	ignored	Already committed, $c(A) = \text{true}$
c2	allowed	
r5(A)	allowed	$c(A) = \text{true}$, $RTS(A) > WTS(A)$
c5	allowed	

1.3

request	response of the system	any changes to state for A <i>and/or</i> explanation of why action wasn't accepted
r2(A)	allowed to read A(0)	$RTS(A(0)) = 20$, $WTS(A(0)) = 0$
w4(A)	Allowed to write A(40)	$WTS(A(40)) = 40$, $RTS(A(40)) = 0$
r3(A)	Allowed to read A(0)	Can read from A(0), A(40), A(0) is valid; $RTS(A(0)) = 30$
w1(A)	Denied; rollback	No such value of A where $TS > WTS(A)$.
w2(A)	Denied; rollback	$WTS < TS(A(0))$, every value of the read timestamp is more.
r5(A)	Allowed to read A(40)	$RTS(A(40)) = 50$

Problem 2: Replication and distributed concurrency control

2.1) 11 copies of each item, fully distributed locking

voting scheme	would it work? (yes/no)	explanation
1a) update 6, read 5	no	Although we update a majority of copies, we might need 5 stale ones. r is not $> n - w$
1b) update 5, read 8	no	Not updating a majority of the copies , although the read condition is satisfied
1c) update 8, read 4	yes	Updating majority ($w > n/2$) and also reading $> n-w$ because $4 > 11-8$
1d) update 3, read 10	no	Not updating majority of copies

2.2) 11 copies of each item, primary-copy locking

voting scheme	would it work? (yes/no)	explanation
1a) update 6, read 5	no	r is not $> n-w$ since $5 \neq 11 - 6$
1b) update 5, read 8	yes	$r > n - w \Rightarrow 8 > 11-5=6$, read majority of copies
1c) update 8, read 4	yes	$r > n - w \Rightarrow 4 > 11-8 = 3$, read majority
1d) update 3, read 10	yes	$r > n - w \Rightarrow 10 > 11-3 = 8$, read majority

2.3)

Using a configuration with more reads can guarantee fault tolerance. **1) (d) would be a good choice** because we have **10 reads out of 11, and 3 updates**. In a read-heavy workload, we ensure **not to read any stale data** because $r > n-w$ and even if we have a site go down, we have a **very high probability of reading an updated value without incurring latency costs**. Networking would not be too much of a bottleneck either since the data can be read from 10 locations.

2.4)

For a write heavy workload, although it may feel like **1 (c) is a good choice** because we have 8 writes, writes usually incur a higher latency cost. It would seem wise to use **1 (b)** because you have a majority of writers which can deal with failed machines and still ensure consistency, and at the same time ensure we **do not incur too high of a writing cost**. 1(d) is not the best option either because we have **10 readers and only 3 writers**, which is not very useful in a write-heavy workload.